

**ME759****High Performance Computing for Engineering Applications****Date Assigned: October 14, 2013****Date Due: October 21, 2013****Cutoff Time: 11:59 PM**

In this problem, you will have to produce a version V5 and a version V6 of the 1D Stencil code that we used in class. To this end, start with the **testV4.cu** version provided in HW directory.

**testV5.cu:** Use shared memory to speed up your execution.


**testV6.cu:** Builds on top of V5 and reduces the run time by considering pinned host memory transactions.

What you will have to deliver:

a) Run a scaling analysis using  $N=103, 104, 105, \dots, 108$  elements and generate a **png** plot that shows GPU-V5 performance against CPU performance. Upload this plot onto the Forum.


Answer:

Plot: (LOG SCALE):



b) The same as above, but shows GPU-V6 performance against CPU performance


Answer: Plot: (LOG SCALE)



c) Generate a png plot that shows the GPU-V5 performance against GPU-V6 performance.

Answer:

Plot: (LINEAR SCALE)



d) What change has had more impact? Why is that the case?

**Answer:**

The change after pinning the memory on the host has more impact than 'weights' and 'in' being stored in the Shared Memory. Pinned memory is memory allocated using the *cudaMallocHost* function, which prevents the memory from being swapped out and provides improved transfer speeds. So, it has more impact than storing 'weights' and 'in' variables in shared memory

**Grading.**

Your submission will be graded as follows:

i) Functionality: 40%

- Program runs on Euler, producing correct results.

ii) Report: 60%

- You provide correct results for a) through d) above.